.

Configuring an Application Load Balancer

Step 1: Create a target group

1.1 Navigate to the AWS console home dashboard, search for and click on **EC2**

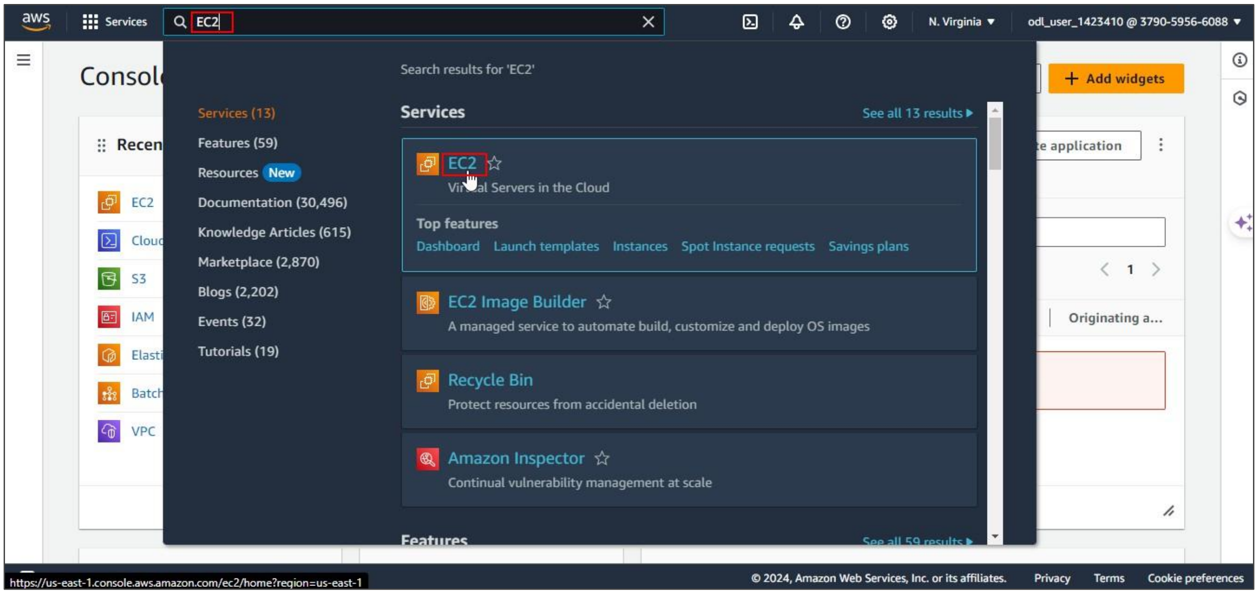

.

1.2 Navigate to the **Load Balancing** section and click on **Target Groups**

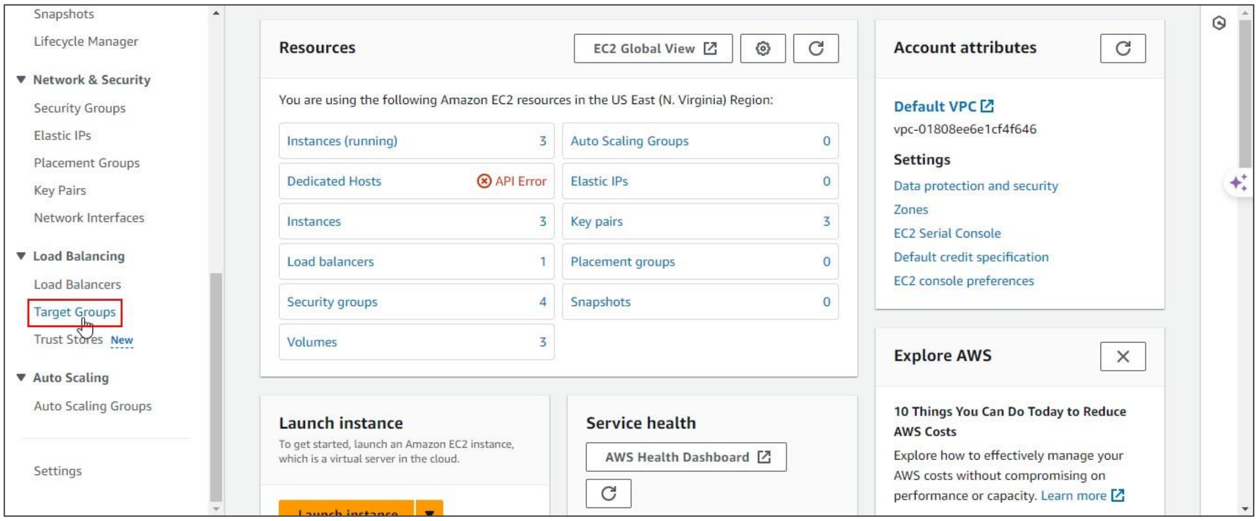

Snapshots
Lifecycle Manager
▼ Network & Security
Security Groups
Elastic IPs
Placement Groups
Key Pairs
Network Interfaces
▼ Load Balancing
Load Balancers
Target Groups
Trust Stores New
▼ Auto Scaling
Auto Scaling Groups
Settings

Resources EC2 Global View

You are using the following Amazon EC2 resources in the US East (N. Virginia) Region:

Instances (running)	3	Auto Scaling Groups	0
Dedicated Hosts	⊗ API Error	Elastic IPs	0
Instances	3	Key pairs	3
Load balancers	1	Placement groups	0
Security groups	4	Snapshots	0
Volumes	3		

Account attributes

Default VPC
vpc-01808ee6e1cf4f646

Settings
Data protection and security
Zones
EC2 Serial Console
Default credit specification
EC2 console preferences

Explore AWS

10 Things You Can Do Today to Reduce AWS Costs
Explore how to effectively manage your AWS costs without compromising on performance or capacity. Learn more

Launch instance
To get started, launch an Amazon EC2 instance, which is a virtual server in the cloud.
Launch instance

Service health
AWS Health Dashboard

1.3 Click on **Create target group**

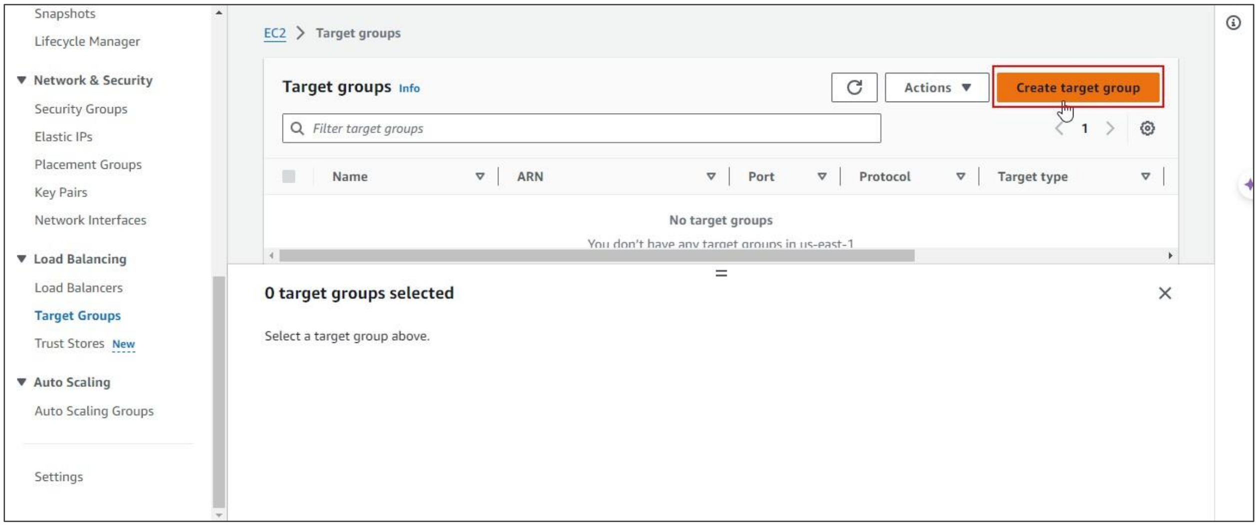

EC2 > Target groups

Target groups Info

Actions ▼ Create target group

Filter target groups

Name	ARN	Port	Protocol	Target type

No target groups
You don't have any target groups in us-east-1

0 target groups selected

Select a target group above.

Snapshots
Lifecycle Manager
▼ Network & Security
Security Groups
Elastic IPs
Placement Groups
Key Pairs
Network Interfaces
▼ Load Balancing
Load Balancers
Target Groups
Trust Stores New
▼ Auto Scaling
Auto Scaling Groups
Settings

.

1.4 In the **Basic configuration** section, choose **Instances** as the target type and enter a name for the target group, such as **MyTargetGroup**

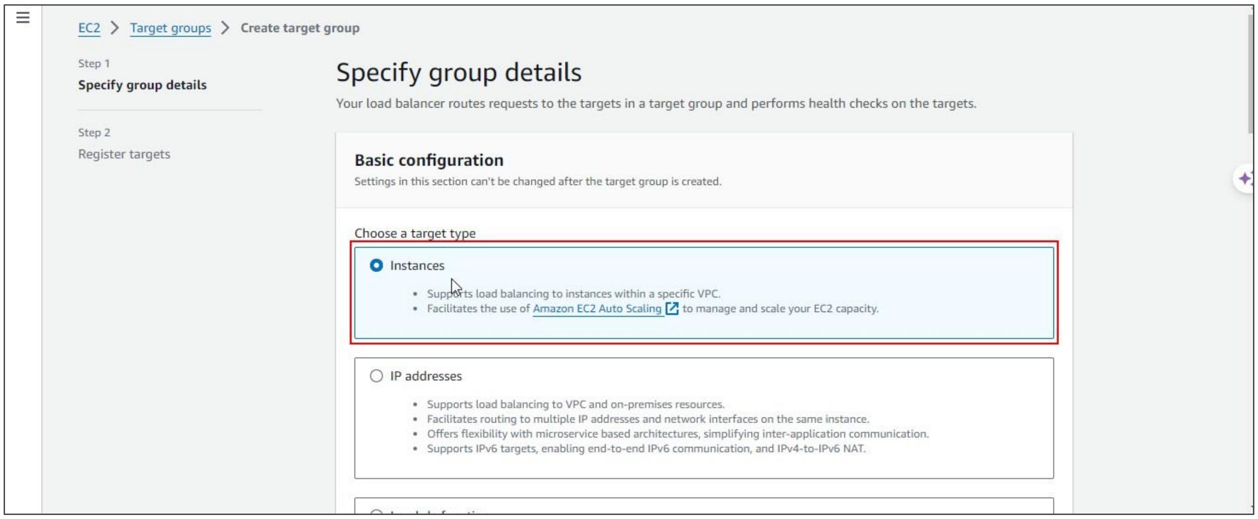

EC2 > Target groups > Create target group

Step 1
Specify group details

Step 2
Register targets

Specify group details

Your load balancer routes requests to the targets in a target group and performs health checks on the targets.

Basic configuration

Settings in this section can't be changed after the target group is created.

Choose a target type

● Instances

- Supports load balancing to instances within a specific VPC.
- Facilitates the use of Amazon EC2 Auto Scaling to manage and scale your EC2 capacity.

○ IP addresses

- Supports load balancing to VPC and on-premises resources.
- Facilitates routing to multiple IP addresses and network interfaces on the same instance.
- Offers flexibility with microservice based architectures, simplifying inter-application communication.
- Supports IPv6 targets, enabling end-to-end IPv6 communication, and IPv4-to-IPv6 NAT.

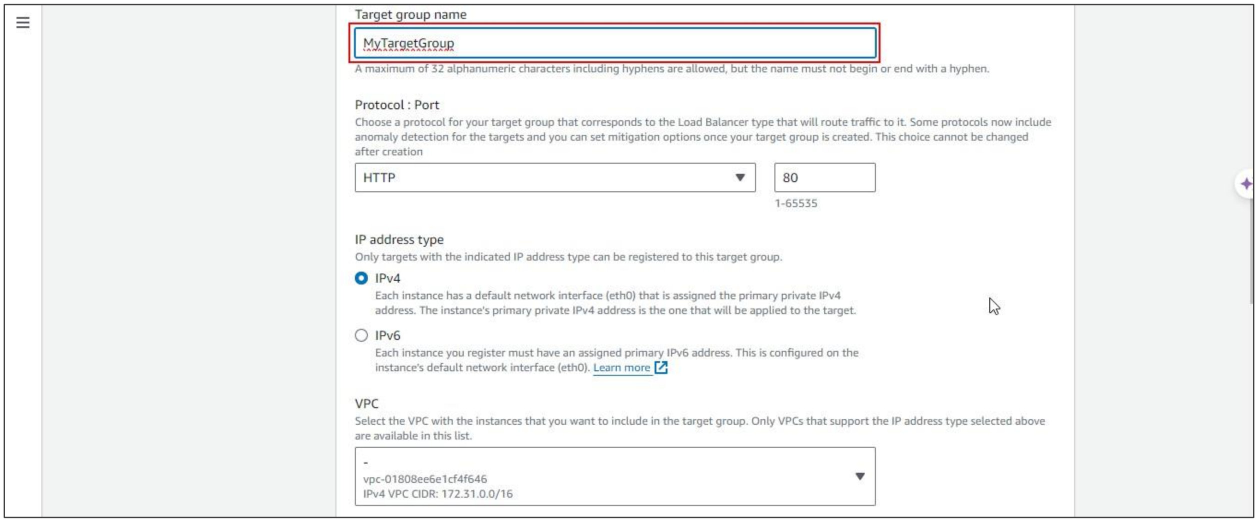

Target group name

MyTargetGroup

A maximum of 32 alphanumeric characters including hyphens are allowed, but the name must not begin or end with a hyphen.

Protocol : Port

Choose a protocol for your target group that corresponds to the Load Balancer type that will route traffic to it. Some protocols now include anomaly detection for the targets and you can set mitigation options once your target group is created. This choice cannot be changed after creation

HTTP ▼ 80

1-65535

IP address type

Only targets with the indicated IP address type can be registered to this target group.

● IPv4
Each instance has a default network interface (eth0) that is assigned the primary private IPv4 address. The instance's primary private IPv4 address is the one that will be applied to the target.

○ IPv6
Each instance you register must have an assigned primary IPv6 address. This is configured on the instance's default network interface (eth0). Learn more

VPC

Select the VPC with the instances that you want to include in the target group. Only VPCs that support the IP address type selected above are available in this list.

-
vpc-01808ee6e1cf4f646
IPv4 VPC CIDR: 172.31.0.0/16 ▼

.

1.5 Set the protocol to **HTTP** and the path to **/index.html** in the **Health checks** section

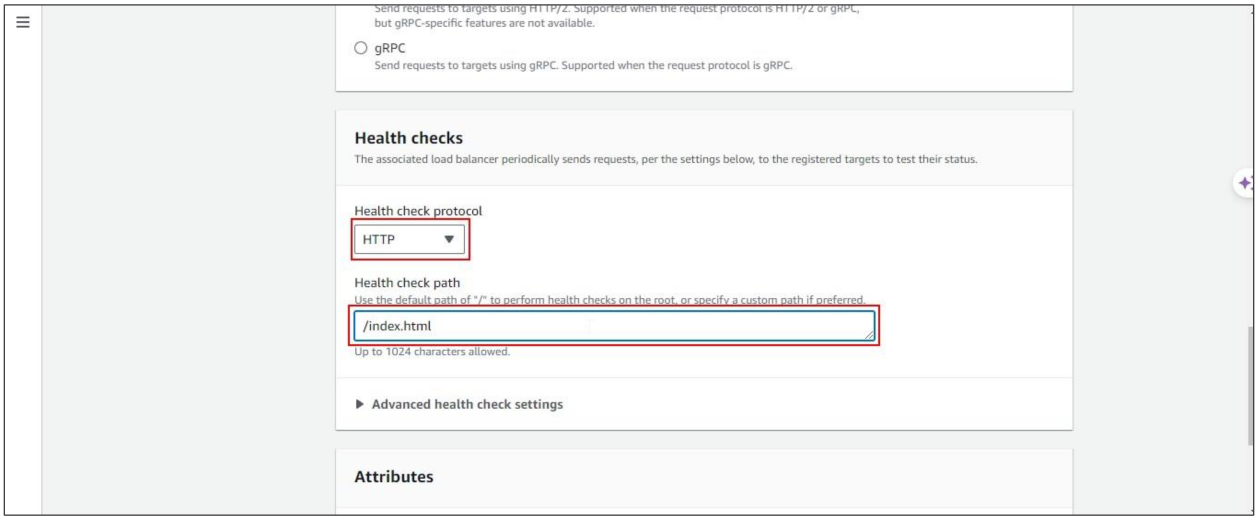

Send requests to targets using HTTP/2. Supported when the request protocol is HTTP/2 or gRPC, but gRPC-specific features are not available.

○ gRPC
Send requests to targets using gRPC. Supported when the request protocol is gRPC.

Health checks
The associated load balancer periodically sends requests, per the settings below, to the registered targets to test their status.

Health check protocol
HTTP ▼

Health check path
Use the default path of "/" to perform health checks on the root, or specify a custom path if preferred.
/index.html
Up to 1024 characters allowed.

▶ Advanced health check settings

Attributes

1.6 Click on **Next**

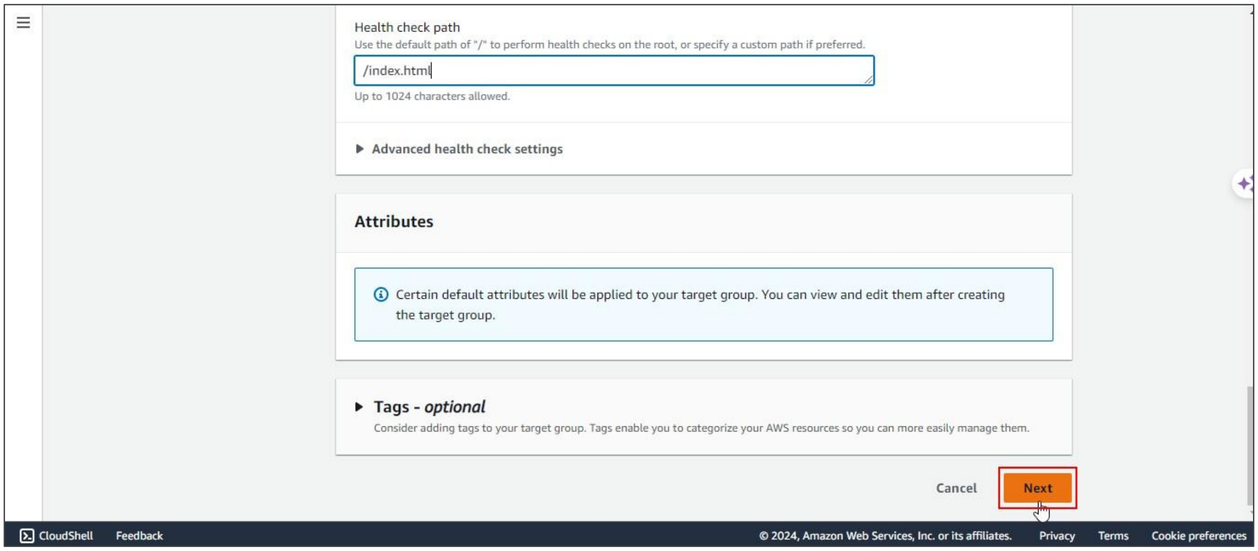

Health check path
Use the default path of "/" to perform health checks on the root, or specify a custom path if preferred.
/index.html
Up to 1024 characters allowed.

▶ Advanced health check settings

Attributes

ⓘ Certain default attributes will be applied to your target group. You can view and edit them after creating the target group.

▶ **Tags - *optional***
Consider adding tags to your target group. Tags enable you to categorize your AWS resources so you can more easily manage them.

.

1.7 Review the configurations and click on **Create target group**

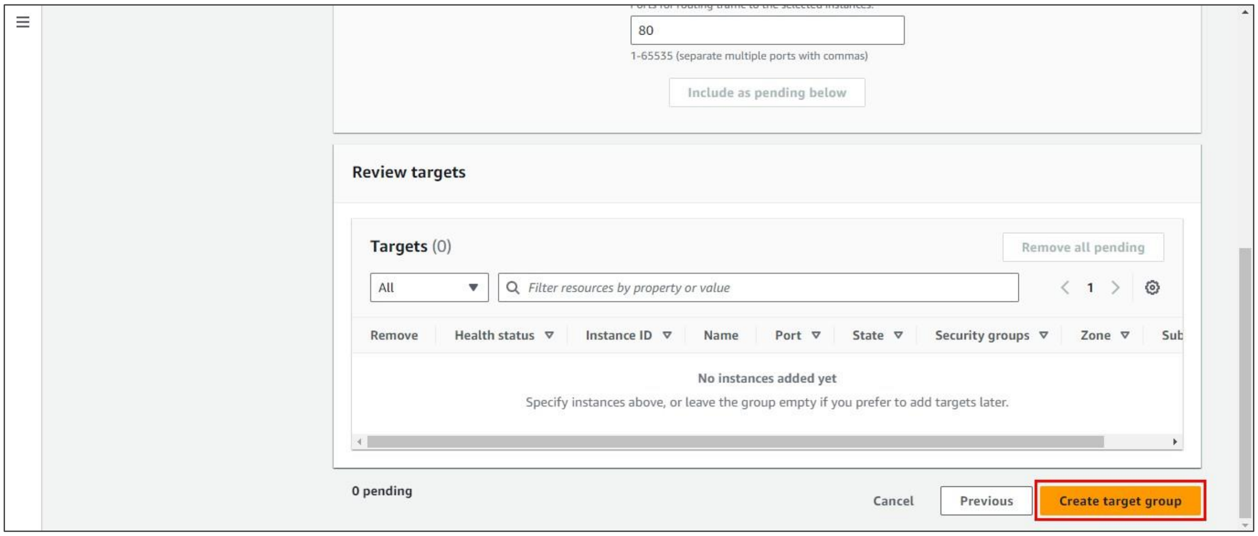

80

1-65535 (separate multiple ports with commas)

Review targets

Targets (0)

Remove	Health status	Instance ID	Name	Port	State	Security groups	Zone	Sub

No instances added yet

Specify instances above, or leave the group empty if you prefer to add targets later.

0 pending

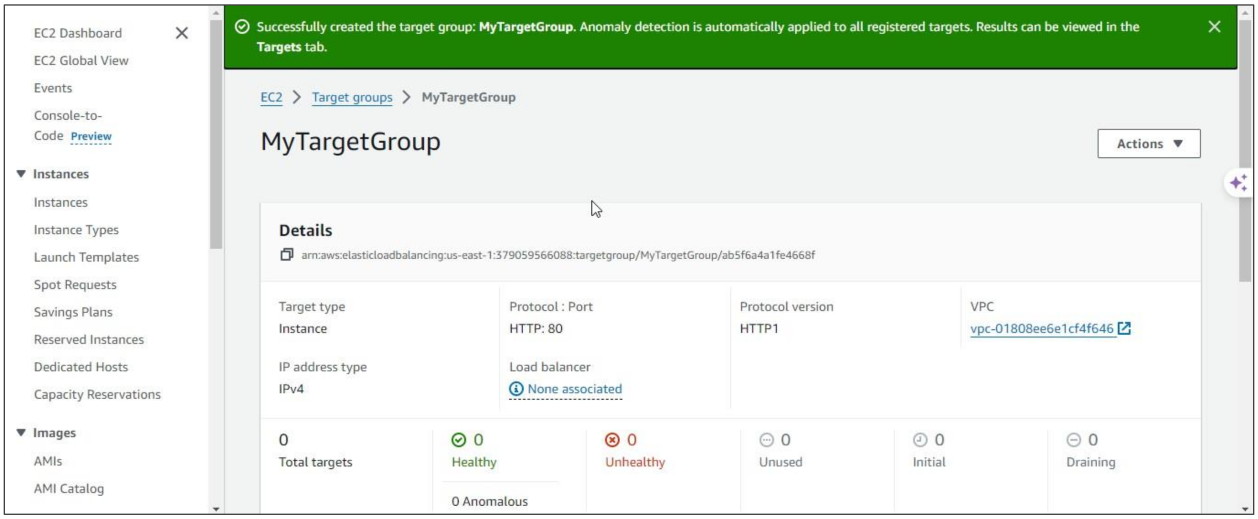

Successfully created the target group: **MyTargetGroup**. Anomaly detection is automatically applied to all registered targets. Results can be viewed in the **Targets** tab.

EC2 > Target groups > MyTargetGroup

MyTargetGroup

Details

arn:aws:elasticloadbalancing:us-east-1:379059566088:targetgroup/MyTargetGroup/ab5f6a4a1fe4668f

Target type	Protocol : Port	Protocol version	VPC
Instance	HTTP: 80	HTTP1	vpc-01808ee6e1cf4f646
IP address type	**Load balancer**		
IPv4	None associated		

0	0	0	0	0	0
Total targets	Healthy	Unhealthy	Unused	Initial	Draining
	0 Anomalous				

The target group has been successfully created.

.

Step 2: Launch EC2 instances

2.1 Navigate to the **Instances** section and click on **Launch instances**

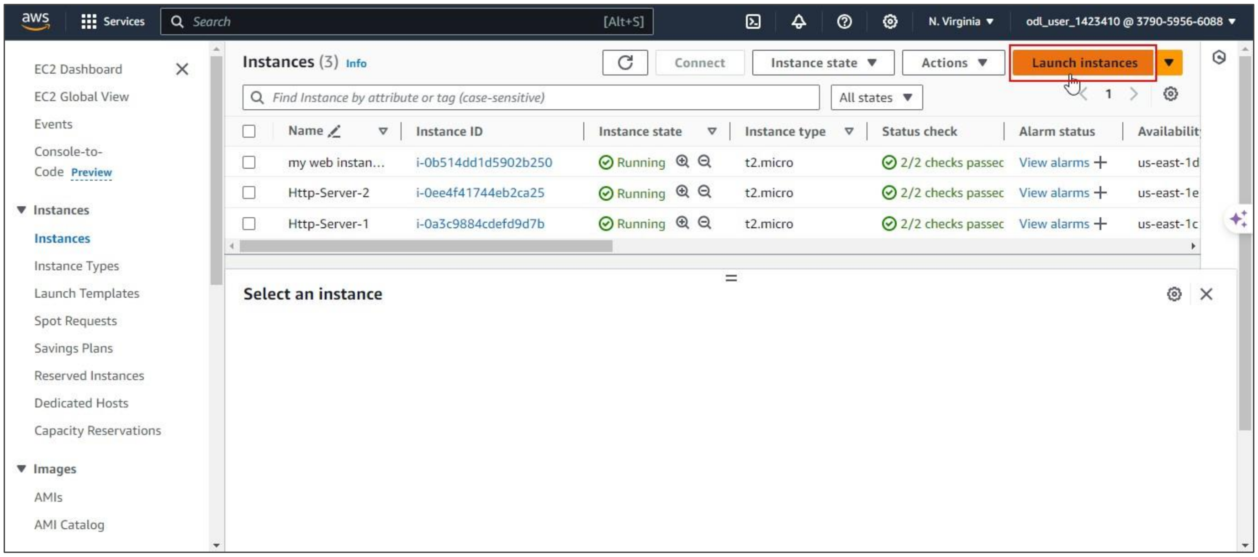

Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availabilit
my web instan…	i-0b514dd1d5902b250	Running	t2.micro	2/2 checks passec	View alarms +	us-east-1d
Http-Server-2	i-0ee4f41744eb2ca25	Running	t2.micro	2/2 checks passec	View alarms +	us-east-1e
Http-Server-1	i-0a3c9884cdefd9d7b	Running	t2.micro	2/2 checks passec	View alarms +	us-east-1c

2.2 Provide a name for the instance and choose an appropriate AMI (**Amazon Linux 2**)

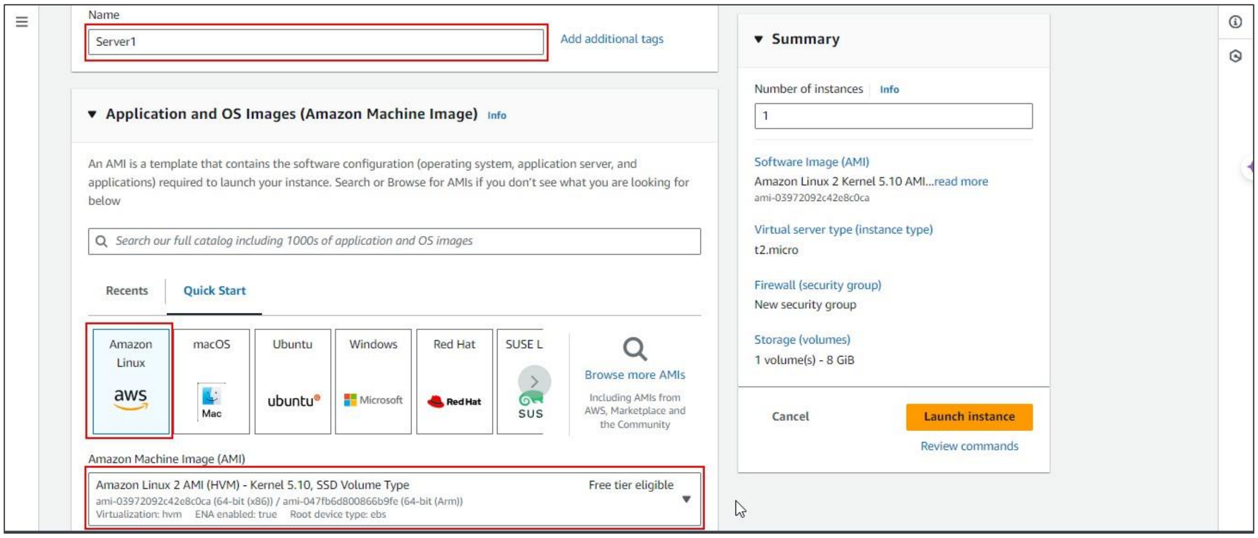
Add additional tags
▼ Application and OS Images (Amazon Machine Image) Info
Search our full catalog including 1000s of application and OS images
Recents
Quick Start
Amazon Linux
aws
macOS
Mac
Ubuntu
ubuntu
Windows
Microsoft
Red Hat
Red Hat
SUSE L
SUS
Browse more AMIs
Including AMIs from AWS, Marketplace and the Community
▼ Summary
Cancel
Launch instance
Review commands

Name: Server1

An AMI is a template that contains the software configuration (operating system, application server, and applications) required to launch your instance. Search or Browse for AMIs if you don't see what you are looking for below

Amazon Machine Image (AMI)
Amazon Linux 2 AMI (HVM) - Kernel 5.10, SSD Volume Type — Free tier eligible
ami-03972092c42e8c0ca (64-bit (x86)) / ami-047fb6d800866b9fe (64-bit (Arm))
Virtualization: hvm ENA enabled: true Root device type: ebs

Number of instances Info: 1

Software Image (AMI): Amazon Linux 2 Kernel 5.10 AMI...read more
ami-03972092c42e8c0ca

Virtual server type (instance type): t2.micro

Firewall (security group): New security group

Storage (volumes): 1 volume(s) - 8 GiB

.

2.3 Select the instance type as **t2.micro**, create a key pair, and name it **Server-1**

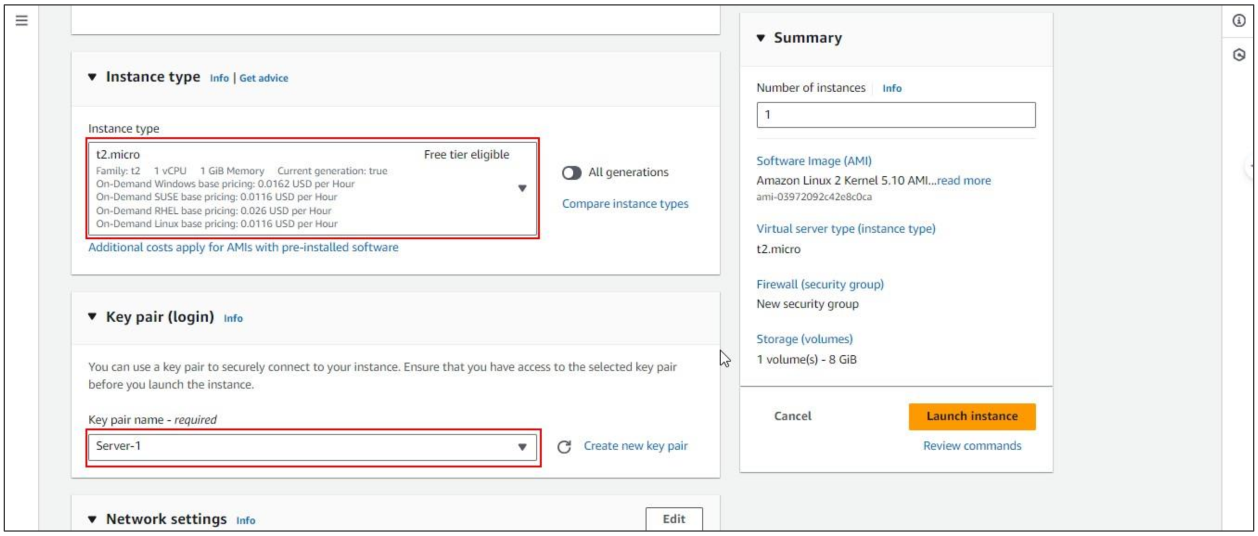

Instance type

Instance type: t2.micro — Free tier eligible

Family: t2 1 vCPU 1 GiB Memory Current generation: true
On-Demand Windows base pricing: 0.0162 USD per Hour
On-Demand SUSE base pricing: 0.0116 USD per Hour
On-Demand RHEL base pricing: 0.026 USD per Hour
On-Demand Linux base pricing: 0.0116 USD per Hour

Additional costs apply for AMIs with pre-installed software

Key pair (login)

You can use a key pair to securely connect to your instance. Ensure that you have access to the selected key pair before you launch the instance.

Key pair name - *required*: Server-1

Summary

Number of instances: 1

Software Image (AMI): Amazon Linux 2 Kernel 5.10 AMI...read more
ami-03972092c42e8c0ca

Virtual server type (instance type): t2.micro

Firewall (security group): New security group

Storage (volumes): 1 volume(s) - 8 GiB

2.4 Configure the network settings as shown:

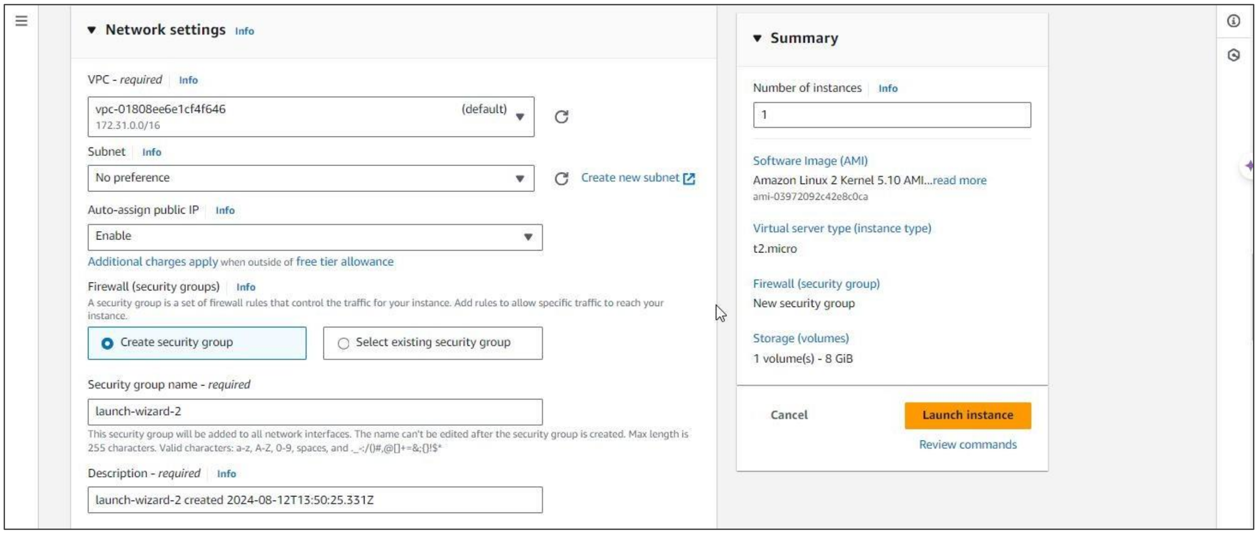

Network settings

VPC - *required*: vpc-01808ee6e1cf4f646 (default)
172.31.0.0/16

Subnet: No preference

Auto-assign public IP: Enable

Additional charges apply when outside of free tier allowance

Firewall (security groups)

A security group is a set of firewall rules that control the traffic for your instance. Add rules to allow specific traffic to reach your instance.

● Create security group ○ Select existing security group

Security group name - *required*: launch-wizard-2

This security group will be added to all network interfaces. The name can't be edited after the security group is created. Max length is 255 characters. Valid characters: a-z, A-Z, 0-9, spaces, and ._-:/()#,@[]+=&;{}!$*

Description - *required*: launch-wizard-2 created 2024-08-12T13:50:25.331Z

Summary

Number of instances: 1

Software Image (AMI): Amazon Linux 2 Kernel 5.10 AMI...read more
ami-03972092c42e8c0ca

Virtual server type (instance type): t2.micro

Firewall (security group): New security group

Storage (volumes): 1 volume(s) - 8 GiB

.

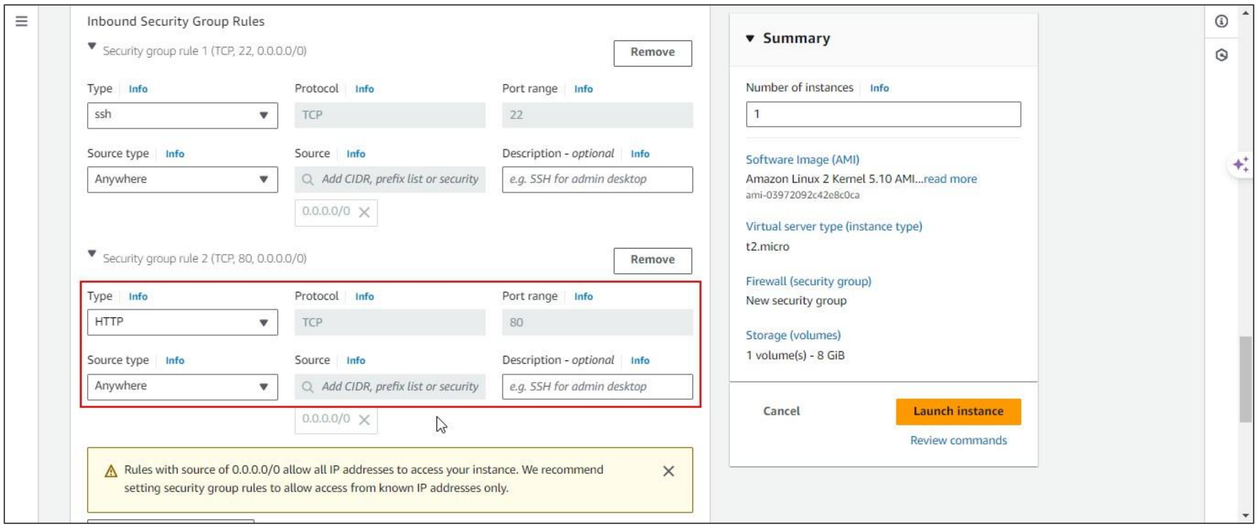

Inbound Security Group Rules

Security group rule 1 (TCP, 22, 0.0.0.0/0)

Type: ssh
Protocol: TCP
Port range: 22
Source type: Anywhere
Source: 0.0.0.0/0
Description - optional: e.g. SSH for admin desktop

Security group rule 2 (TCP, 80, 0.0.0.0/0)

Type: HTTP
Protocol: TCP
Port range: 80
Source type: Anywhere
Source: 0.0.0.0/0
Description - optional: e.g. SSH for admin desktop

Rules with source of 0.0.0.0/0 allow all IP addresses to access your instance. We recommend setting security group rules to allow access from known IP addresses only.

Summary

Number of instances: 1
Software Image (AMI): Amazon Linux 2 Kernel 5.10 AMI...read more ami-03972092c42e8c0ca
Virtual server type (instance type): t2.micro
Firewall (security group): New security group
Storage (volumes): 1 volume(s) - 8 GiB

2.5 Add the following user data script in the **Advance details** section, and click on **Launch instance:**

#!/bin/bash
yum update -y
yum install httpd -y
echo "<html><body><h1>This is Webserver1</h1></body></html>" >
/var/www/html/index.html
systemctl start httpd
systemctl enable httpd

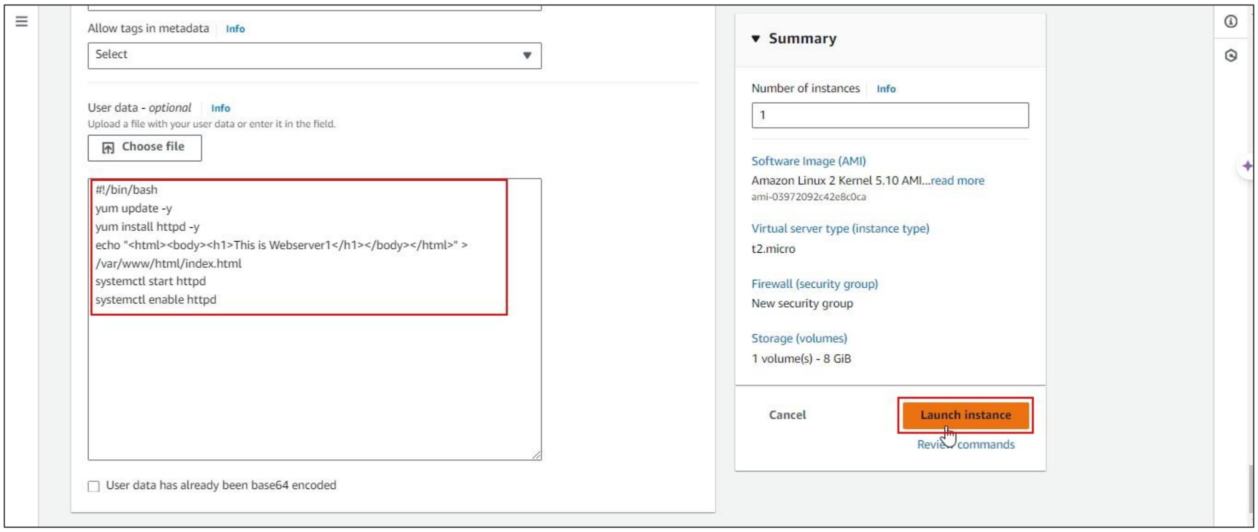

Allow tags in metadata: Select

User data - optional
Upload a file with your user data or enter it in the field.

```
#!/bin/bash
yum update -y
yum install httpd -y
echo "<html><body><h1>This is Webserver1</h1></body></html>" >
/var/www/html/index.html
systemctl start httpd
systemctl enable httpd
```

□ User data has already been base64 encoded

Summary

Number of instances: 1
Software Image (AMI): Amazon Linux 2 Kernel 5.10 AMI...read more ami-03972092c42e8c0ca
Virtual server type (instance type): t2.micro
Firewall (security group): New security group
Storage (volumes): 1 volume(s) - 8 GiB

.

You will see the following interface:

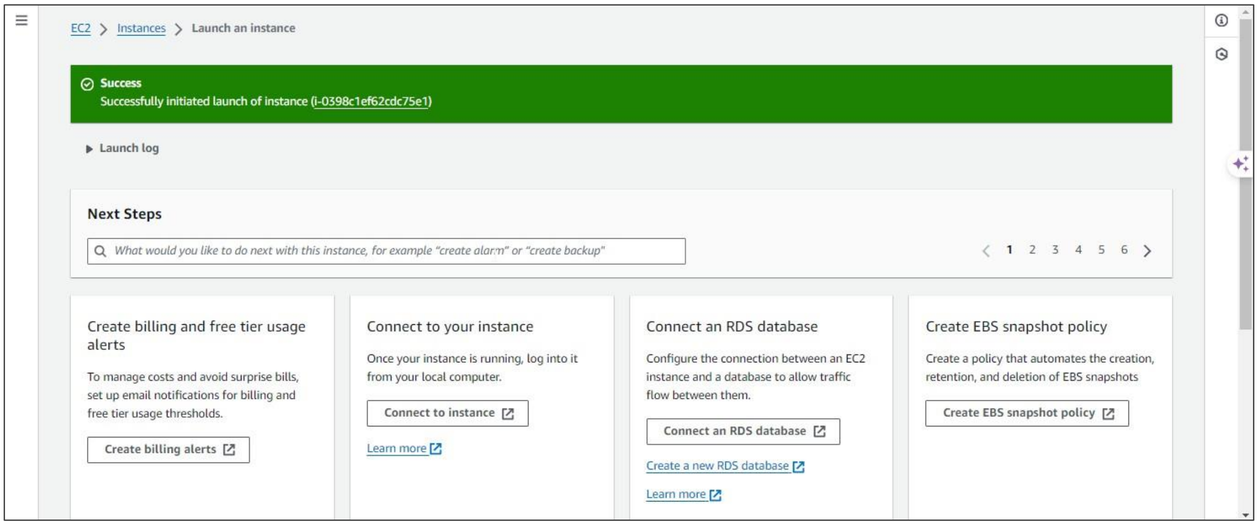

Success
Successfully initiated launch of instance (i-0398c1ef62cdc75e1)

Next Steps

What would you like to do next with this instance, for example "create alarm" or "create backup"

Create billing and free tier usage alerts

To manage costs and avoid surprise bills, set up email notifications for billing and free tier usage thresholds.

Connect to your instance

Once your instance is running, log into it from your local computer.

Connect an RDS database

Configure the connection between an EC2 instance and a database to allow traffic flow between them.

Create EBS snapshot policy

Create a policy that automates the creation, retention, and deletion of EBS snapshots

2.6 Launch another EC2 instance using the same steps, but modify the user data script to display the message **This is Webserver2**

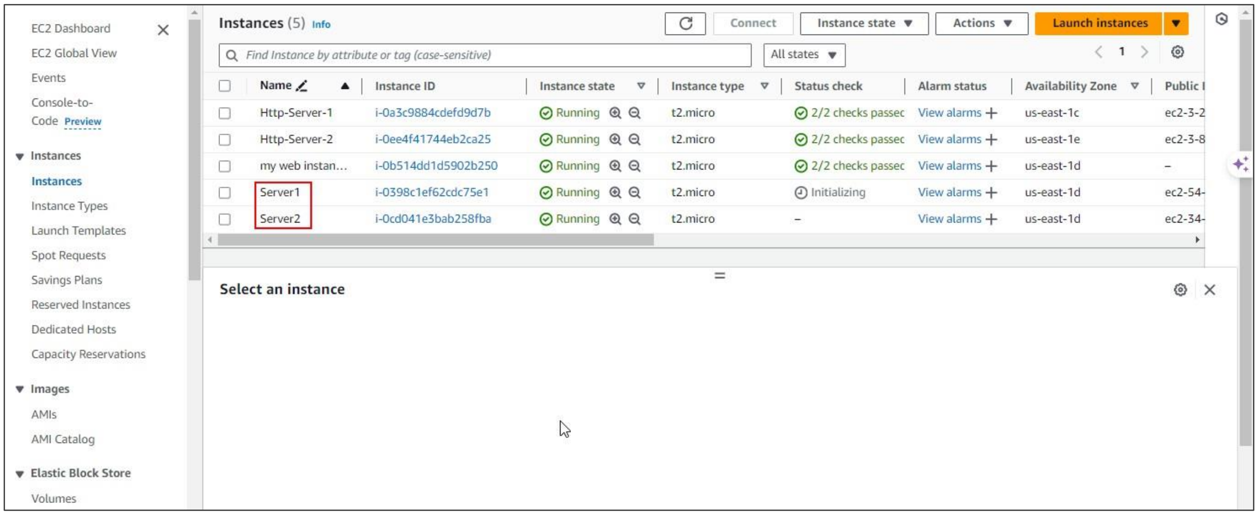

Instances (5) Info

Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone	Public I
Http-Server-1	i-0a3c9884cdefd9d7b	Running	t2.micro	2/2 checks passec	View alarms +	us-east-1c	ec2-3-2
Http-Server-2	i-0ee4f41744eb2ca25	Running	t2.micro	2/2 checks passec	View alarms +	us-east-1e	ec2-3-8
my web instan...	i-0b514dd1d5902b250	Running	t2.micro	2/2 checks passec	View alarms +	us-east-1d	–
Server1	i-0398c1ef62cdc75e1	Running	t2.micro	Initializing	View alarms +	us-east-1d	ec2-54-
Server2	i-0cd041e3bab258fba	Running	t2.micro	–	View alarms +	us-east-1d	ec2-34-

The EC2 instances have been successfully launched.

.

Step 3: Configure the target group

3.1 Navigate to the **Target Groups** section and select the target group created in Step 1

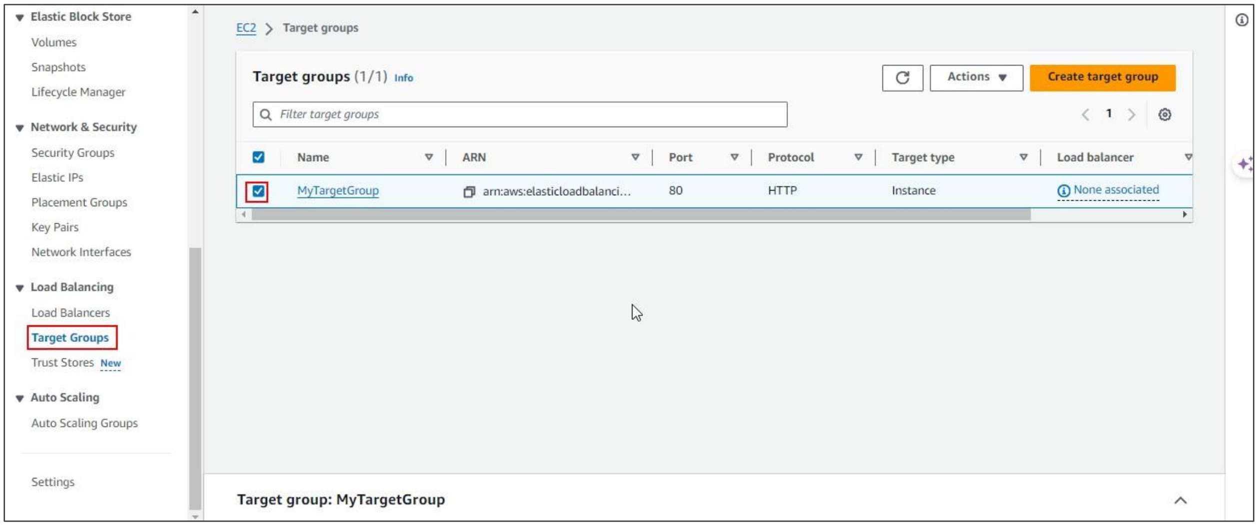

☑	Name	ARN	Port	Protocol	Target type	Load balancer
☑	MyTargetGroup	arn:aws:elasticloadbalanci...	80	HTTP	Instance	None associated

3.2 Click on **Register targets** from the **Actions** menu

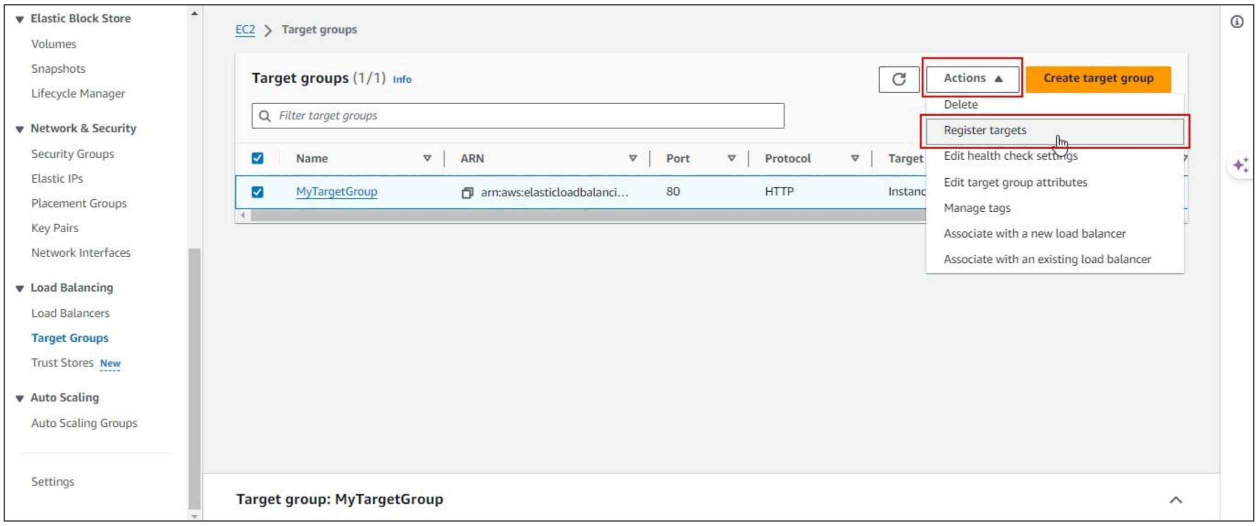

☑	Name	ARN	Port	Protocol	Target
☑	MyTargetGroup	arn:aws:elasticloadbalanci...	80	HTTP	Instanc

3.3 Select the instances (**Server1** and **Server2**) that were launched in Step 2 and click on **Include as pending below**

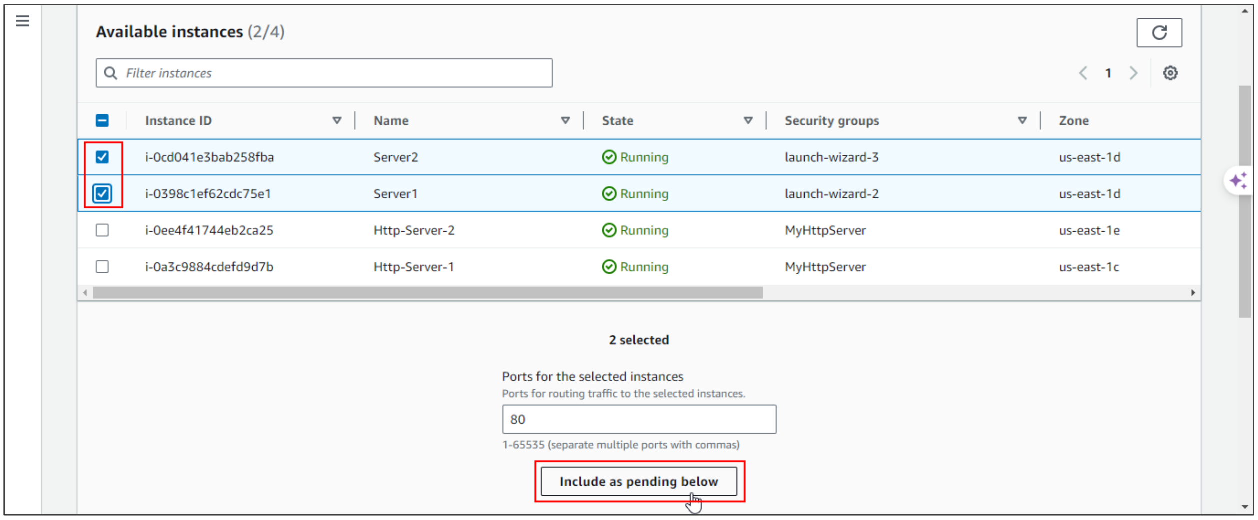

Available instances (2/4)

Filter instances

	Instance ID	Name	State	Security groups	Zone
☑	i-0cd041e3bab258fba	Server2	✓ Running	launch-wizard-3	us-east-1d
☑	i-0398c1ef62cdc75e1	Server1	✓ Running	launch-wizard-2	us-east-1d
□	i-0ee4f41744eb2ca25	Http-Server-2	✓ Running	MyHttpServer	us-east-1e
□	i-0a3c9884cdefd9d7b	Http-Server-1	✓ Running	MyHttpServer	us-east-1c

2 selected

Ports for the selected instances

Ports for routing traffic to the selected instances.

80

1-65535 (separate multiple ports with commas)

3.4 Click on **Register pending targets** to register the instances with the target group

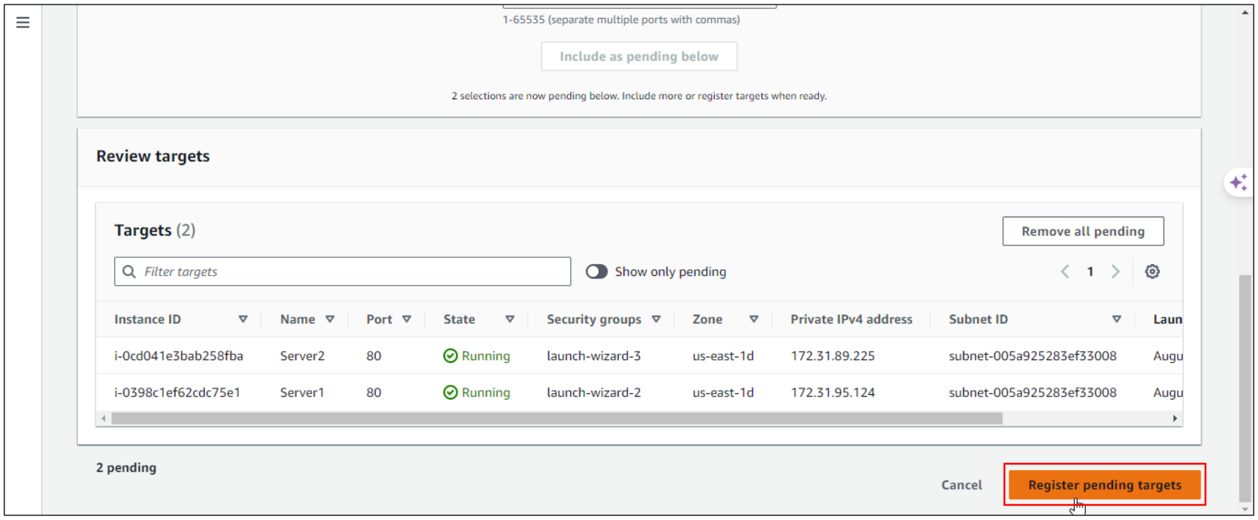

1-65535 (separate multiple ports with commas)

2 selections are now pending below. Include more or register targets when ready.

Review targets

Targets (2)

Instance ID	Name	Port	State	Security groups	Zone	Private IPv4 address	Subnet ID	Laun
i-0cd041e3bab258fba	Server2	80	✓ Running	launch-wizard-3	us-east-1d	172.31.89.225	subnet-005a925283ef33008	Augu
i-0398c1ef62cdc75e1	Server1	80	✓ Running	launch-wizard-2	us-east-1d	172.31.95.124	subnet-005a925283ef33008	Augu

2 pending

.

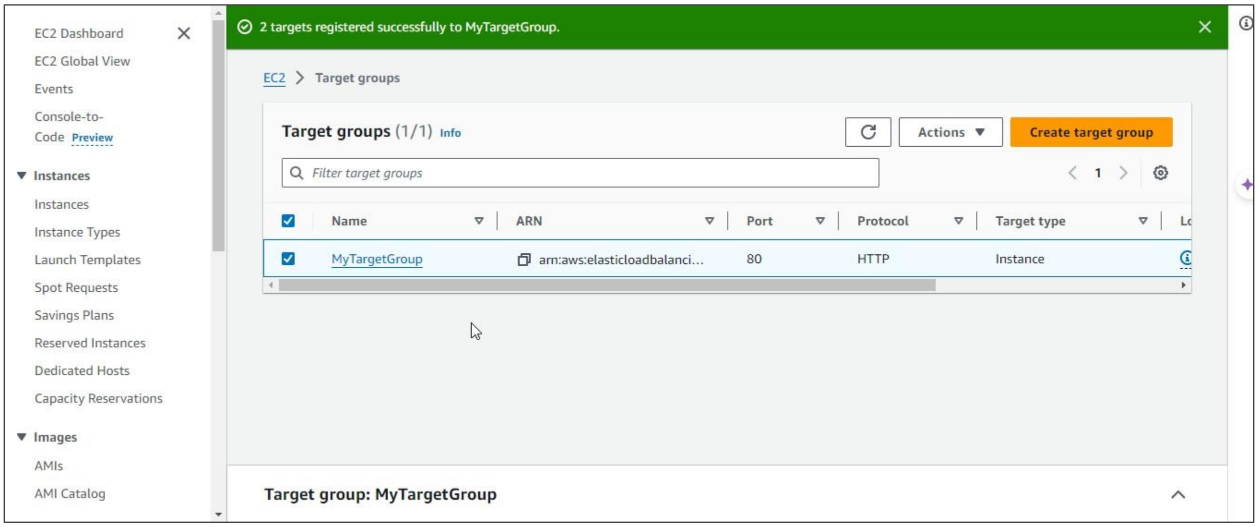

2 targets registered successfully to MyTargetGroup.

Target groups (1/1) Info

☑	Name	ARN	Port	Protocol	Target type
☑	MyTargetGroup	arn:aws:elasticloadbalanci...	80	HTTP	Instance

Target group: MyTargetGroup

Step 4: Create a Load Balancer

4.1 Navigate to the **Load Balancers** section under **Load Balancing** and click **Create load balancer**

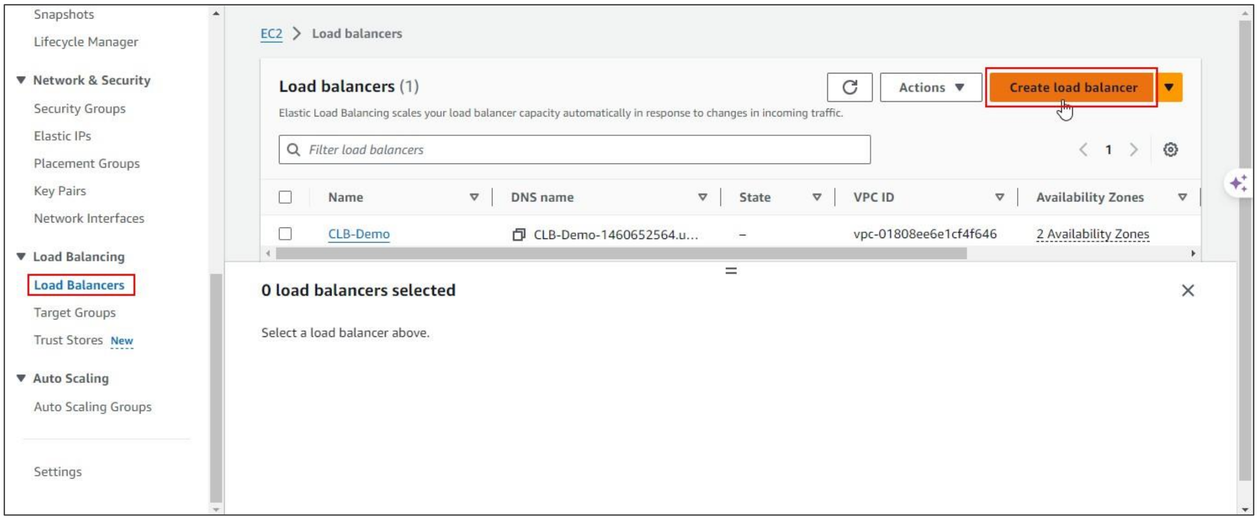

Load balancers (1)

Elastic Load Balancing scales your load balancer capacity automatically in response to changes in incoming traffic.

□	Name	DNS name	State	VPC ID	Availability Zones
□	CLB-Demo	CLB-Demo-1460652564.u...	–	vpc-01808ee6e1cf4f646	2 Availability Zones

0 load balancers selected

Select a load balancer above.

.

4.2 Choose **Application Load Balancer** and click **Create**

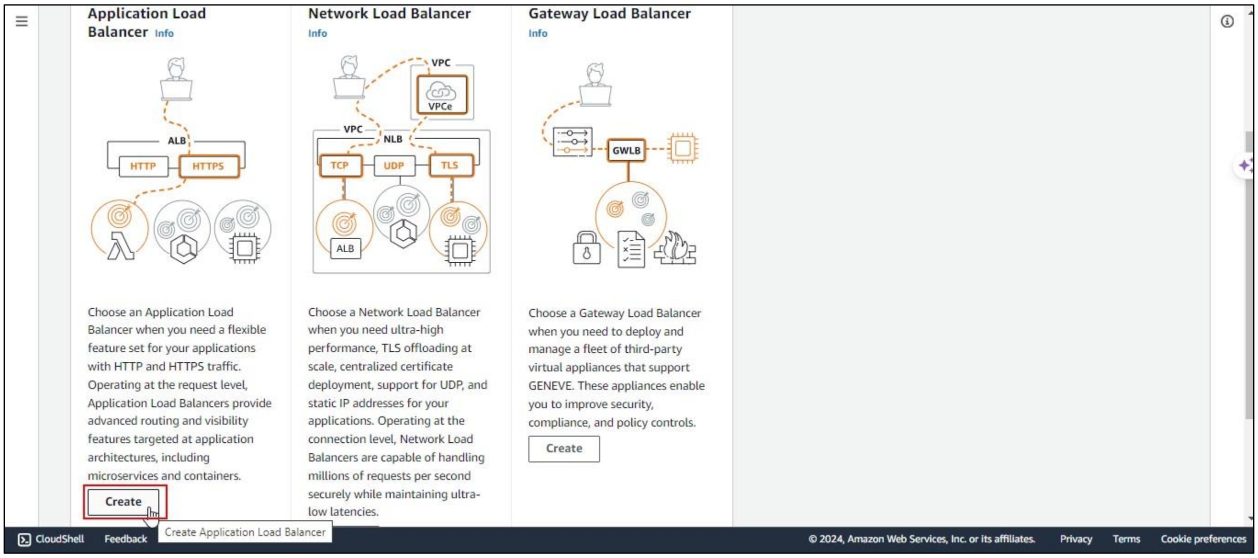

Application Load Balancer Info

Choose an Application Load Balancer when you need a flexible feature set for your applications with HTTP and HTTPS traffic. Operating at the request level, Application Load Balancers provide advanced routing and visibility features targeted at application architectures, including microservices and containers.

Create

Network Load Balancer Info

Choose a Network Load Balancer when you need ultra-high performance, TLS offloading at scale, centralized certificate deployment, support for UDP, and static IP addresses for your applications. Operating at the connection level, Network Load Balancers are capable of handling millions of requests per second securely while maintaining ultra-low latencies.

Gateway Load Balancer Info

Choose a Gateway Load Balancer when you need to deploy and manage a fleet of third-party virtual appliances that support GENEVE. These appliances enable you to improve security, compliance, and policy controls.

Create

CloudShell
Feedback
Create Application Load Balancer
© 2024, Amazon Web Services, Inc. or its affiliates.
Privacy
Terms
Cookie preferences

4.3 Configure the load balancer settings, enter a name for the load balancer, such as **my-alb**, and select availability zones, such as **us-east-1a** and **us-east-1b**

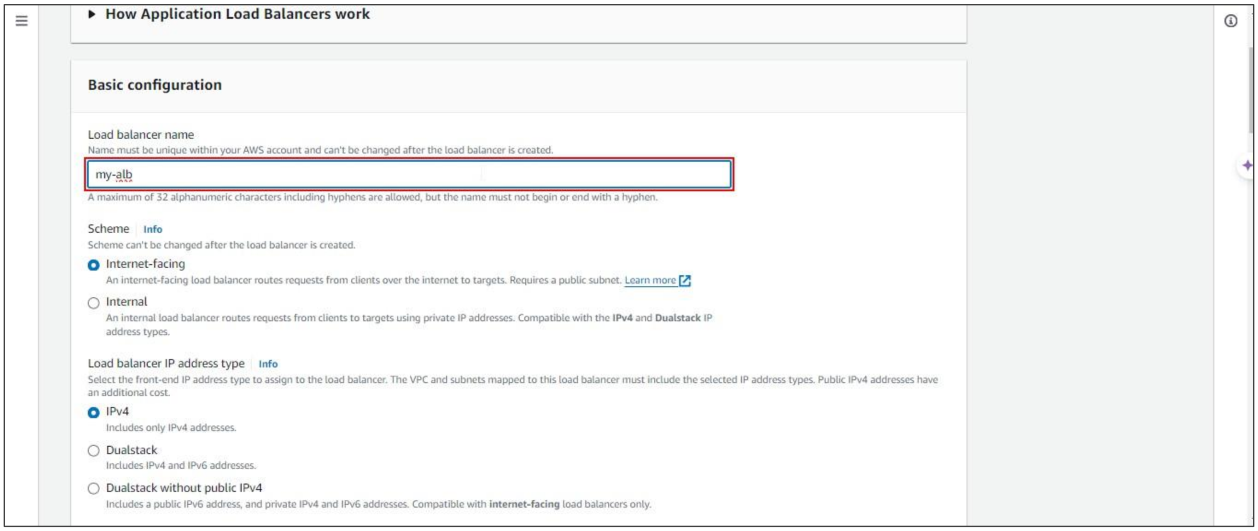

▸ How Application Load Balancers work

Basic configuration

Load balancer name
Name must be unique within your AWS account and can't be changed after the load balancer is created.
my-alb
A maximum of 32 alphanumeric characters including hyphens are allowed, but the name must not begin or end with a hyphen.

Scheme | Info
Scheme can't be changed after the load balancer is created.

● Internet-facing
An internet-facing load balancer routes requests from clients over the internet to targets. Requires a public subnet. Learn more

○ Internal
An internal load balancer routes requests from clients to targets using private IP addresses. Compatible with the **IPv4** and **Dualstack** IP address types.

Load balancer IP address type | Info
Select the front-end IP address type to assign to the load balancer. The VPC and subnets mapped to this load balancer must include the selected IP address types. Public IPv4 addresses have an additional cost.

● IPv4
Includes only IPv4 addresses.

○ Dualstack
Includes IPv4 and IPv6 addresses.

○ Dualstack without public IPv4
Includes a public IPv6 address, and private IPv4 and IPv6 addresses. Compatible with **internet-facing** load balancers only.

.

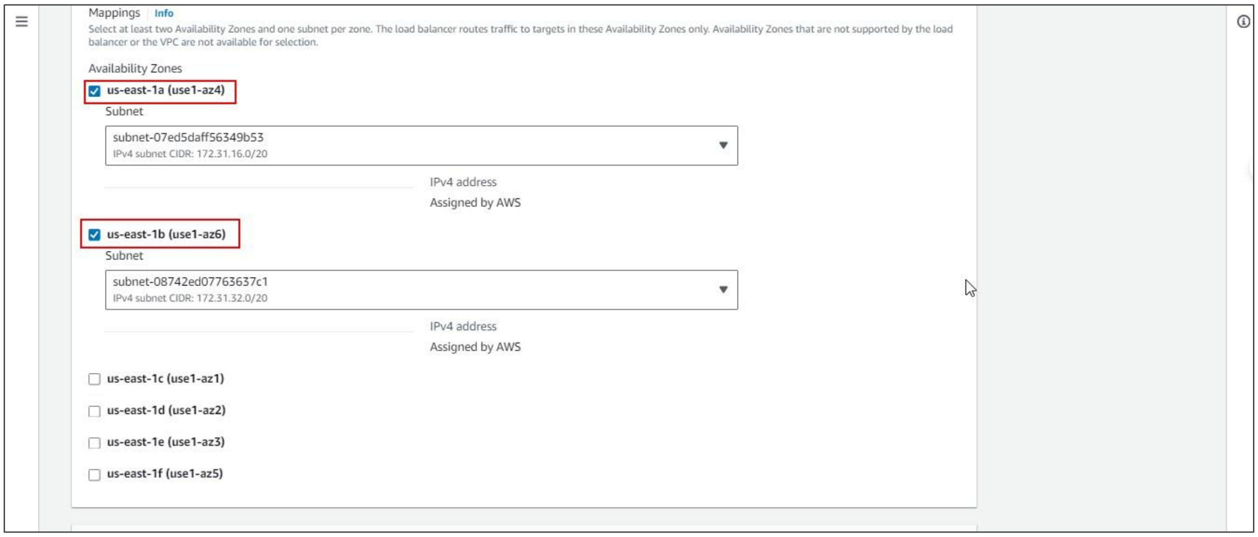

Mappings Info

Select at least two Availability Zones and one subnet per zone. The load balancer routes traffic to targets in these Availability Zones only. Availability Zones that are not supported by the load balancer or the VPC are not available for selection.

Availability Zones

☑ us-east-1a (use1-az4)

Subnet

subnet-07ed5daff56349b53
IPv4 subnet CIDR: 172.31.16.0/20

IPv4 address
Assigned by AWS

☑ us-east-1b (use1-az6)

Subnet

subnet-08742ed07763637c1
IPv4 subnet CIDR: 172.31.32.0/20

IPv4 address
Assigned by AWS

□ us-east-1c (use1-az1)

□ us-east-1d (use1-az2)

□ us-east-1e (use1-az3)

□ us-east-1f (use1-az5)

4.4 Choose the default action for the listener configuration to accept HTTP traffic on port **80**, and select the target group created in Step 1

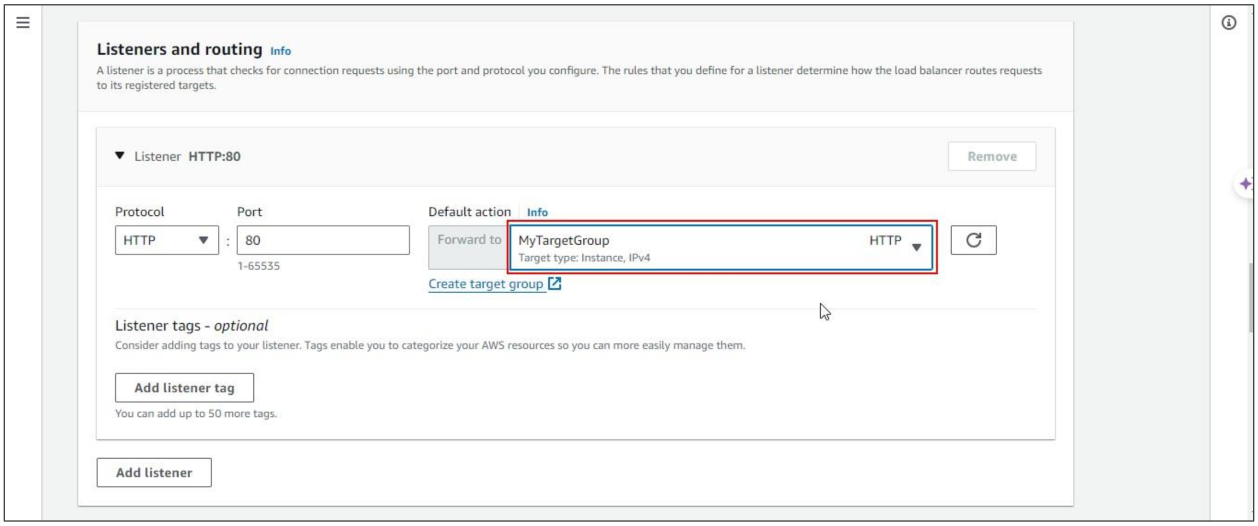

Listeners and routing Info

A listener is a process that checks for connection requests using the port and protocol you configure. The rules that you define for a listener determine how the load balancer routes requests to its registered targets.

▼ Listener **HTTP:80** Remove

Protocol: HTTP

Port: 80
1-65535

Default action Info

Forward to: MyTargetGroup HTTP
Target type: Instance, IPv4

Create target group

Listener tags - *optional*

Consider adding tags to your listener. Tags enable you to categorize your AWS resources so you can more easily manage them.

Add listener tag

You can add up to 50 more tags.

Add listener

.

4.5 Review the configuration and click **Create load balancer**

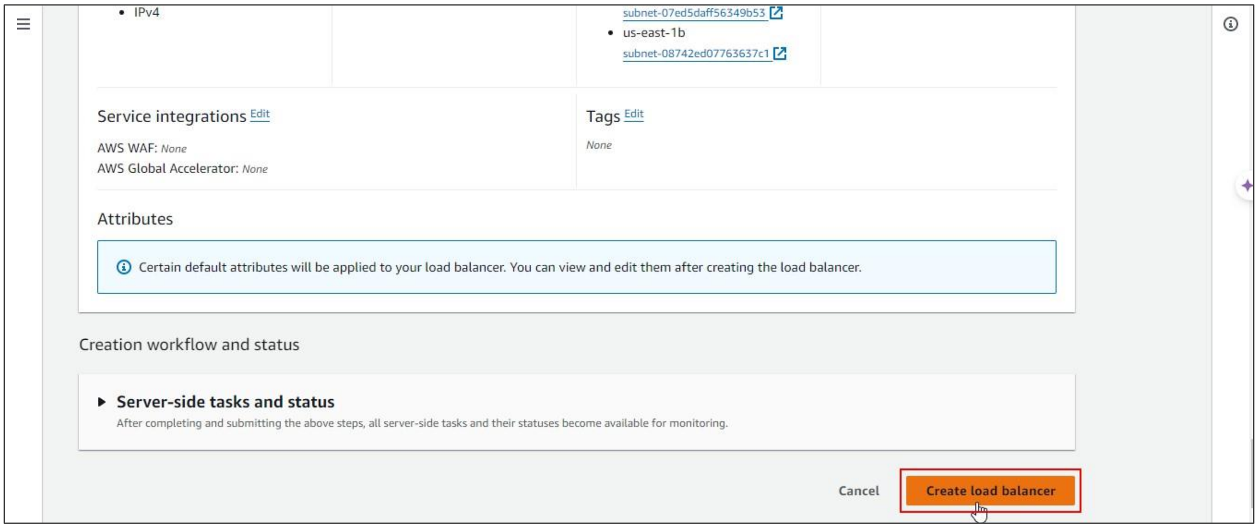

- IPv4

subnet-07ed5daff56349b53

- us-east-1b

subnet-08742ed07763637c1

Service integrations Edit

AWS WAF: *None*

AWS Global Accelerator: *None*

Tags Edit

None

Attributes

Certain default attributes will be applied to your load balancer. You can view and edit them after creating the load balancer.

Creation workflow and status

Server-side tasks and status

After completing and submitting the above steps, all server-side tasks and their statuses become available for monitoring.

Wait until the **Status** changes from **Provisioning** to **Active**

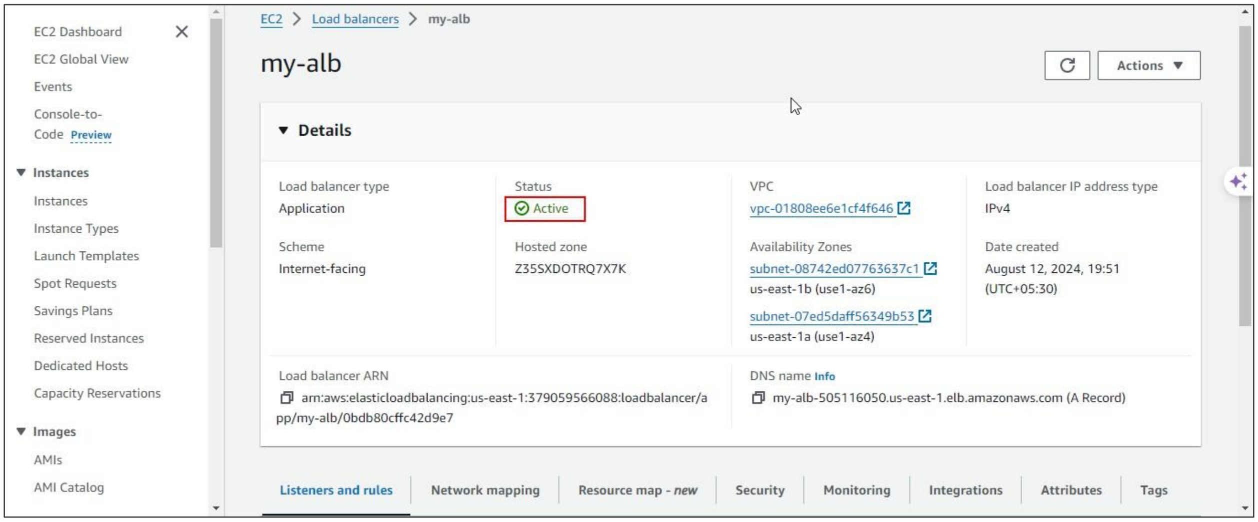

my-alb

Details

Load balancer type: Application

Status: Active

VPC: vpc-01808ee6e1cf4f646

Load balancer IP address type: IPv4

Scheme: Internet-facing

Hosted zone: Z35SXDOTRQ7X7K

Availability Zones: subnet-08742ed07763637c1 us-east-1b (use1-az6); subnet-07ed5daff56349b53 us-east-1a (use1-az4)

Date created: August 12, 2024, 19:51 (UTC+05:30)

Load balancer ARN: arn:aws:elasticloadbalancing:us-east-1:379059566088:loadbalancer/app/my-alb/0bdb80cffc42d9e7

DNS name Info: my-alb-505116050.us-east-1.elb.amazonaws.com (A Record)

.

4.6 Click on the **Security** tab

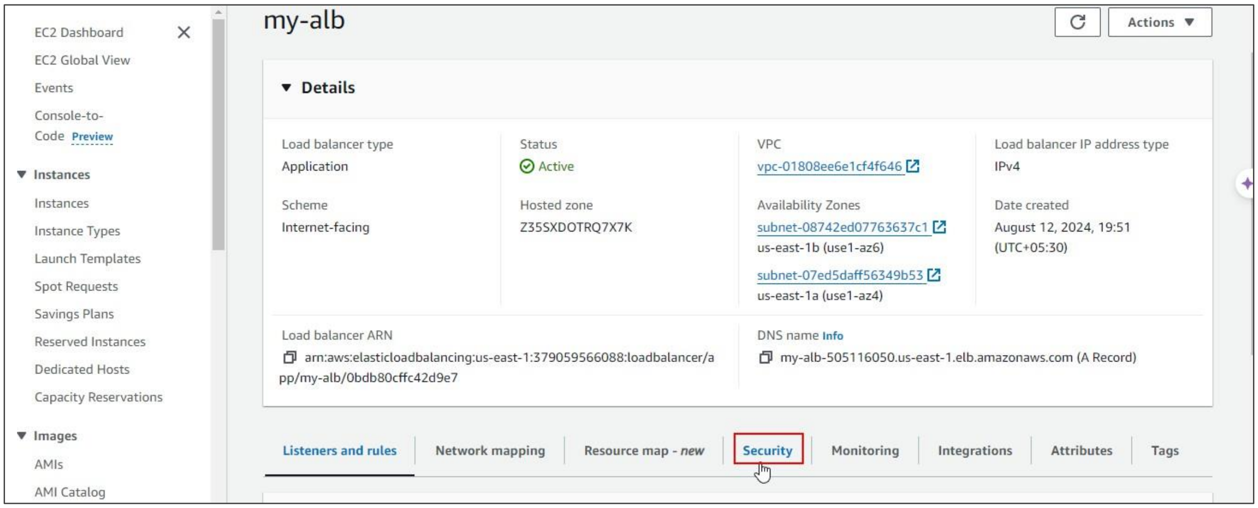

my-alb

Details

Load balancer type: Application

Status: Active

VPC: vpc-01808ee6e1cf4f646

Load balancer IP address type: IPv4

Scheme: Internet-facing

Hosted zone: Z35SXDOTRQ7X7K

Availability Zones: subnet-08742ed07763637c1 us-east-1b (use1-az6); subnet-07ed5daff56349b53 us-east-1a (use1-az4)

Date created: August 12, 2024, 19:51 (UTC+05:30)

Load balancer ARN: arn:aws:elasticloadbalancing:us-east-1:379059566088:loadbalancer/app/my-alb/0bdb80cffc42d9e7

DNS name Info: my-alb-505116050.us-east-1.elb.amazonaws.com (A Record)

4.7 Click on the **Security Group ID** name

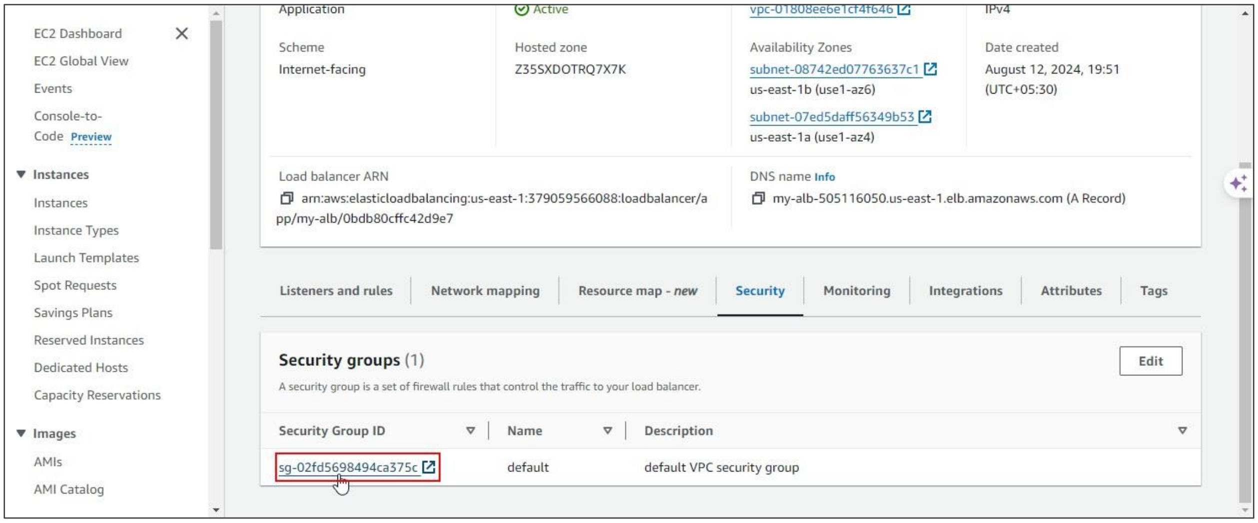

Scheme: Internet-facing

Hosted zone: Z35SXDOTRQ7X7K

Availability Zones: subnet-08742ed07763637c1 us-east-1b (use1-az6); subnet-07ed5daff56349b53 us-east-1a (use1-az4)

Date created: August 12, 2024, 19:51 (UTC+05:30)

Load balancer ARN: arn:aws:elasticloadbalancing:us-east-1:379059566088:loadbalancer/app/my-alb/0bdb80cffc42d9e7

DNS name Info: my-alb-505116050.us-east-1.elb.amazonaws.com (A Record)

Security groups (1)

A security group is a set of firewall rules that control the traffic to your load balancer.

Security Group ID	Name	Description
sg-02fd5698494ca375c	default	default VPC security group

4.8 Click on **Edit inbound rules**

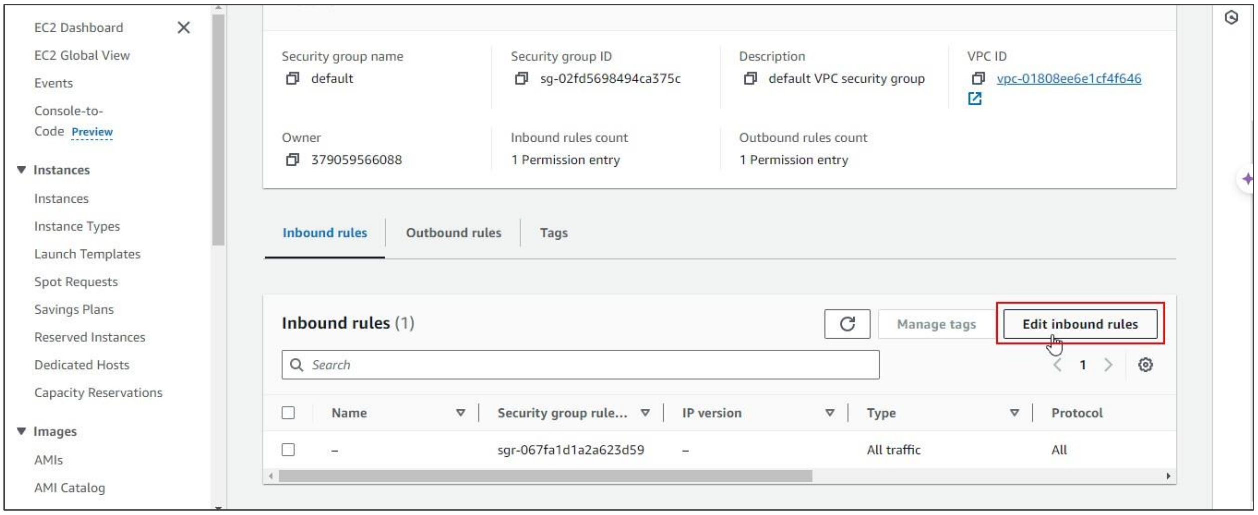

Security group name: default

Security group ID: sg-02fd5698494ca375c

Description: default VPC security group

VPC ID: vpc-01808ee6e1cf4f646

Owner: 379059566088

Inbound rules count: 1 Permission entry

Outbound rules count: 1 Permission entry

Inbound rules (1)

	Name	Security group rule...	IP version	Type	Protocol
□	–	sgr-067fa1d1a2a623d59	–	All traffic	All

4.9 Create an inbound rule to permit port **80** access for all, and click on **Save rules** as shown:

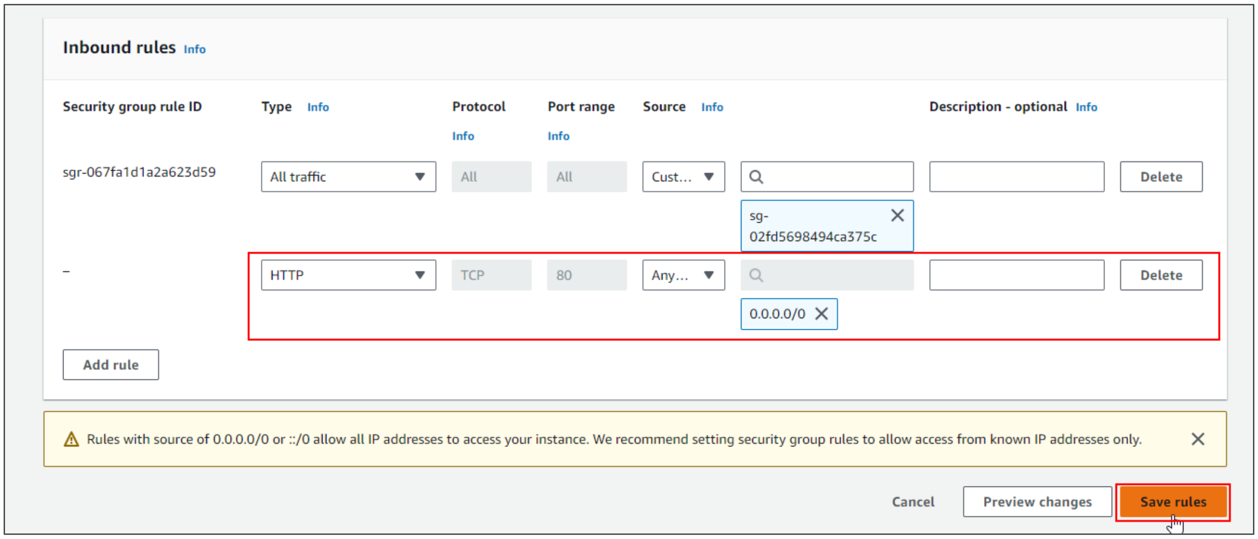
Inbound rules Info

Security group rule ID	Type Info	Protocol Info	Port range Info	Source Info		Description - optional Info	
sgr-067fa1d1a2a623d59	All traffic	All	All	Cust...	sg-02fd5698494ca375c		Delete
–	HTTP	TCP	80	Any...	0.0.0.0/0		Delete

⚠ Rules with source of 0.0.0.0/0 or ::/0 allow all IP addresses to access your instance. We recommend setting security group rules to allow access from known IP addresses only.

.

Step 5: Test the Load Balancer

1.1 Navigate to the **Target Groups** section and select the target group you created

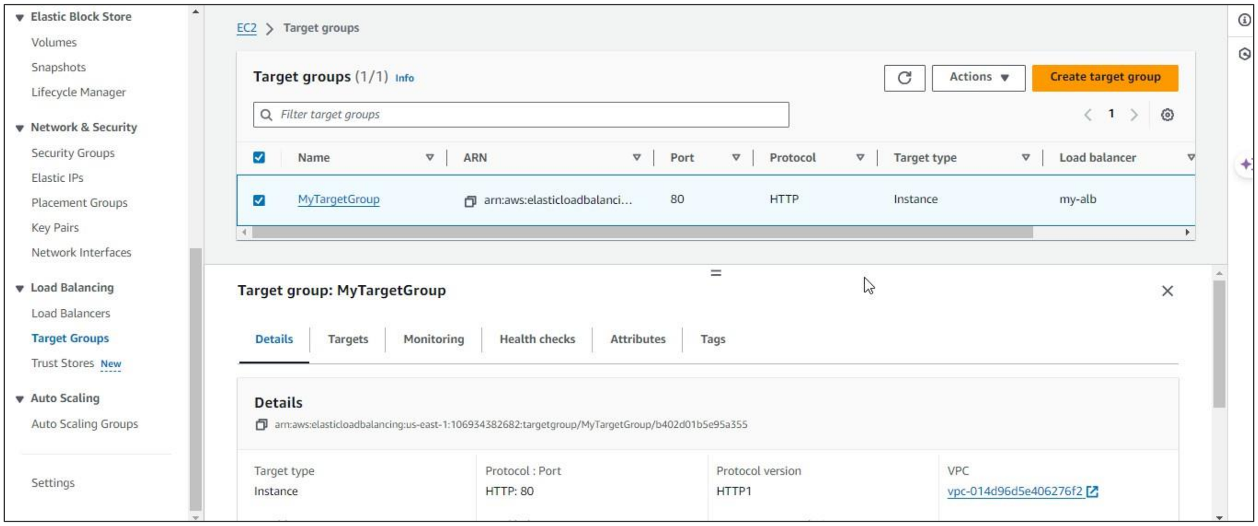

Target groups (1/1) Info

☑	Name	ARN	Port	Protocol	Target type	Load balancer
☑	MyTargetGroup	arn:aws:elasticloadbalanci...	80	HTTP	Instance	my-alb

Target group: MyTargetGroup

Details

arn:aws:elasticloadbalancing:us-east-1:106934382682:targetgroup/MyTargetGroup/b402d01b5e95a355

Target type	Protocol : Port	Protocol version	VPC
Instance	HTTP: 80	HTTP1	vpc-014d96d5e406276f2

1.2 Click on **Details** to verify that your instances are registered and healthy

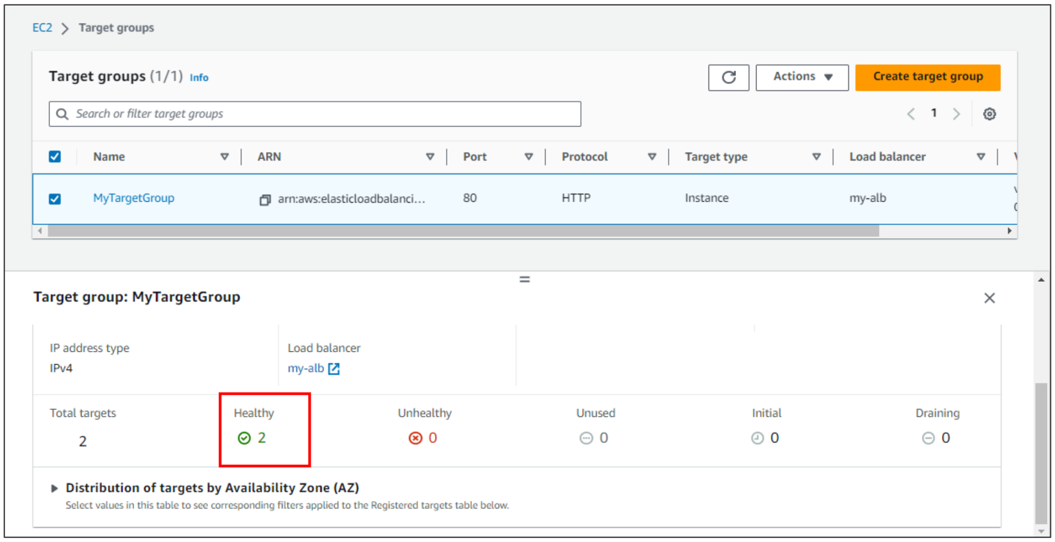

Target groups (1/1) Info

☑	Name	ARN	Port	Protocol	Target type	Load balancer
☑	MyTargetGroup	arn:aws:elasticloadbalanci...	80	HTTP	Instance	my-alb

Target group: MyTargetGroup

IP address type	Load balancer
IPv4	my-alb

Total targets	Healthy	Unhealthy	Unused	Initial	Draining
2	2	0	0	0	0

▶ Distribution of targets by Availability Zone (AZ)

Select values in this table to see corresponding filters applied to the Registered targets table below.

.

1.3 Navigate to the **Load Balancers** section and copy the DNS name of the Load Balancer

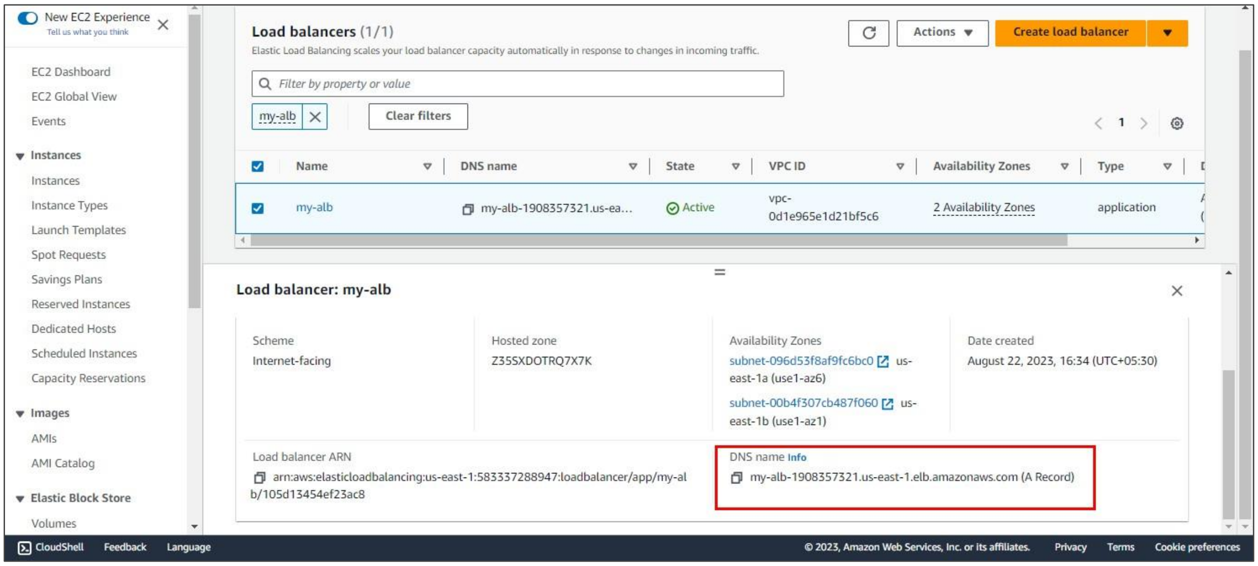

Load balancers (1/1)

Elastic Load Balancing scales your load balancer capacity automatically in response to changes in incoming traffic.

Filter by property or value

my-alb

☑	Name	DNS name	State	VPC ID	Availability Zones	Type
☑	my-alb	my-alb-1908357321.us-ea...	Active	vpc-0d1e965e1d21bf5c6	2 Availability Zones	application

Load balancer: my-alb

Scheme
Internet-facing

Hosted zone
Z35SXDOTRQ7X7K

Availability Zones
subnet-096d53f8af9fc6bc0 us-east-1a (use1-az6)
subnet-00b4f307cb487f060 us-east-1b (use1-az1)

Date created
August 22, 2023, 16:34 (UTC+05:30)

Load balancer ARN
arn:aws:elasticloadbalancing:us-east-1:583337288947:loadbalancer/app/my-alb/105d13454ef23ac8

DNS name Info
my-alb-1908357321.us-east-1.elb.amazonaws.com (A Record)

© 2023, Amazon Web Services, Inc. or its affiliates.

1.4 Open a browser window and paste the DNS URL into the address bar

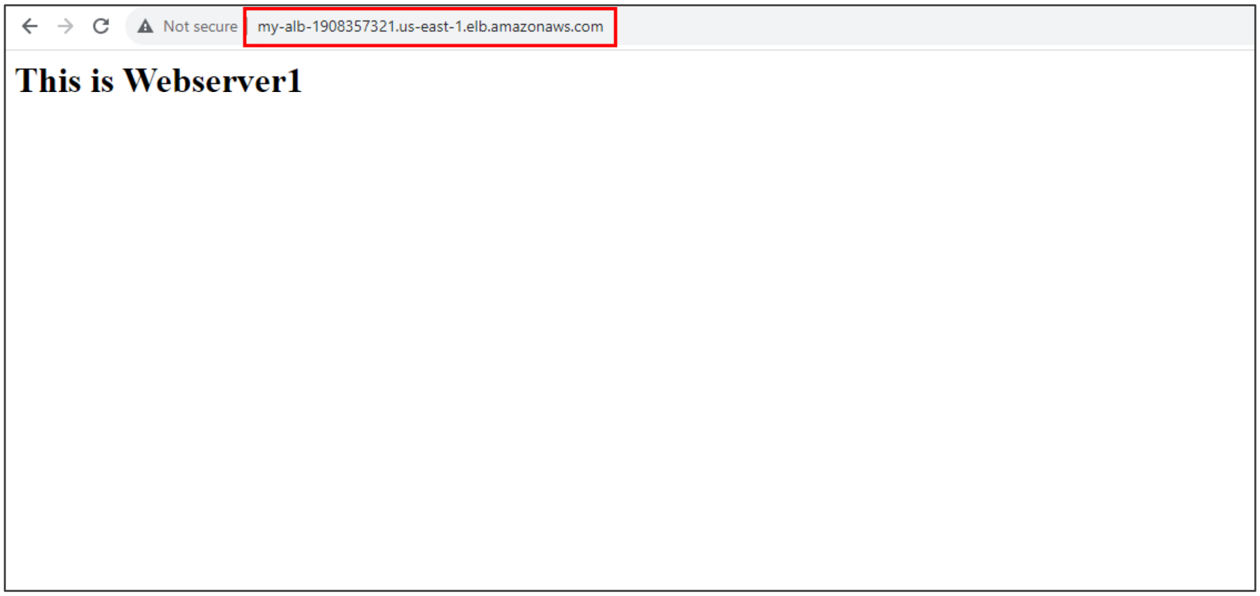

my-alb-1908357321.us-east-1.elb.amazonaws.com

This is Webserver1

You will observe the header message originating from the **Server1** instance.

.

1.5 Refresh the web page multiple times to see the header message originating from the **Server2** instance

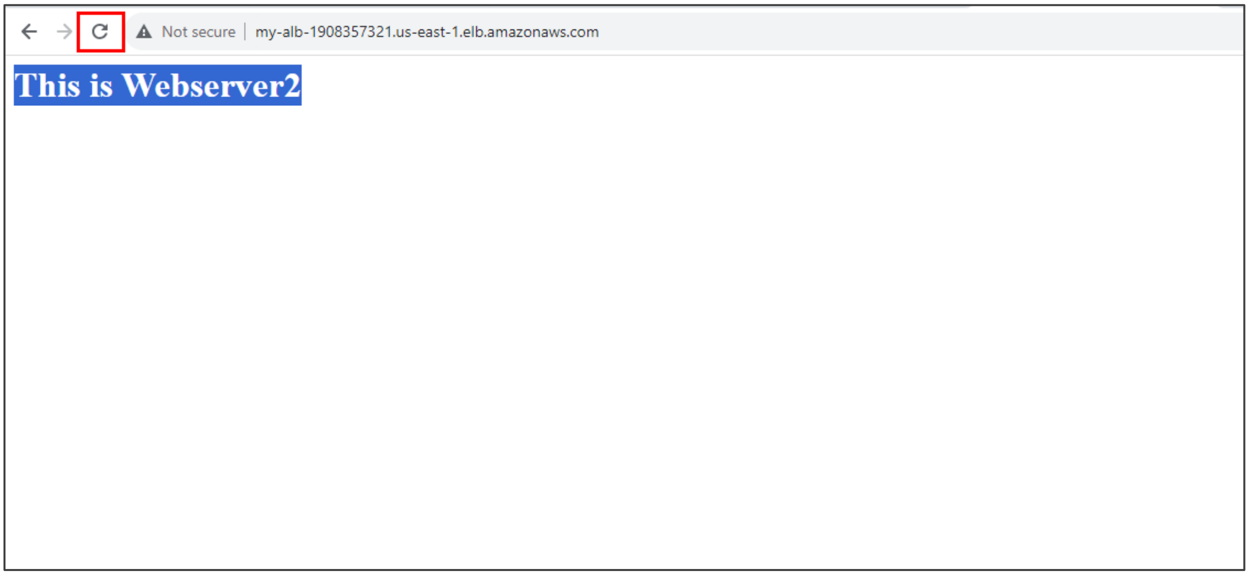

Not secure | my-alb-1908357321.us-east-1.elb.amazonaws.com

This is Webserver2

By following these steps, you have successfully configured an Application Load Balancer in AWS to distribute traffic across multiple EC2 instances, ensuring load balancing and redundancy.